

Index

Sr No.	Title of Experiment	Date	Marks	Signature
1	Study of Deep learning Packages.	08-Aug-22	10	P
2	Implementing Feedforward neural networks with Keras and TensorFlow	22-Aug-22	10	P
3	Build the Image classification model by dividing the model	29-Aug-22	10	P
4	Use Autoencoder to implement anomaly detection.	12-Sep-22	10	P
5	Implement the Continuous Bag of Words (CBOW) Model.	19-Sep-22	10	P
6	Object detection using Transfer Learning of CNN architectures.	03-Oct-22	10	P

Aim: Study of deep learning packaging: Tensorflow, Keras, Theano and PyTorch. Document the distinct features and functionality of the packages.

OBJECTIVE: Introduction to various deep learning tools and how to use them.

Software used: Windows/Linux/Mac, Tensorflow, Keras, PyTorch

THEORY

Tensorflow

- Tensorflow is an open source software library released in 2015 by Google to make it easier for developers to design, build, and train deep learning models. Tensorflow originated as an internal library that Google developers used to build models in-house, and we expect additional functionality to be added to the open source version as it is tested and vetted in the internal flavour.

On high level, Tensorflow is a Python library that allows users to express arbitrary computation as a graph of data flows. Nodes in the graph represent mathematical operations, whereas edges represent data that is communicated from one node to another. Data in Tensorflow is represented as tensors, which are multi-dimensional arrays.

Running Tensorflow

```
import tensorflow as tf
```

```
deep-learning = tf.constant('DeepLearning')
```

```
session = tf.Session()
```


designed. Two (deep learning)

$$a = f_j(\text{constant}(2))$$

$$a = f_j(\text{constant}(2))$$

$$\text{multiply} = f_j \text{ mul}(a, b)$$

used in, sum (multiplying)

TensorFlow Operations

Category

Examples

Element-wise mathematics: Add, sub, mul, Div, Exp, Log, Grate, ReLU, etc.

Array operations: Concat, Slice, Split, Concat, Reshape, etc.

Matrix operations: MatMul, MatrixInverse, Matrix Determinant

Unstacked operations: Variable, Assign, AssignAdd, ...

Neural network building blocks: Softmax, sigmoid, ReLU, Convolution 2D, etc.

Checkpointing operations: Save, Restore

Resource management operations: Engine, Device, AutoAssign, etc.

Control flow operations: Merge, Switch, End, Loop, Nest, etc.

Keras

Keras is an open source deep learning framework for Python.

It has been developed by an AI researcher at Google named François Chollet.

Keras is based on minimalist architecture that provides a clean and easy way to create deep learning models based on TensorFlow or Theano.

Keras is designed to quickly define deep learning models. It is

Features

Keras leverages various optimization techniques to make highly

neural network MLP easier, and more performant. It supports

Conclusion

Keras is a high-level neural network API, written in Python and capable of running on top of TensorFlow, CNTK, or Theano.

It is designed to be easy to use, and to be a good starting point for more complex models.

It is a good choice for those who want to experiment with deep learning.

following features -

- Lightweight, simple, and extensible API.

- Minimal structure - easy to achieve the result without any boilerplate.

- It supports multiple platforms and backends.

- It is user friendly framework which runs on both CPU and GPU.

- Highly stability of computation.

Benefits

Keras is highly powerful and dynamic framework and comes with the following advantages -

- Large community support

- Easy to test.

- Keras neural networks are written in Python which makes things simple.

- Keras support both convolution and recurrent networks.

- Deep learning models are discrete components, so that you can combine them easily.

PyTorch

PyTorch is an open-source deep learning framework for GPUs and CPUs. PyTorch is an open-source deep learning framework based on Python, and Torch is mainly used for applications using GPUs and CPUs. PyTorch is based on the idea of deep learning frameworks like TensorFlow and Keras since it uses dynamic computation graph and is completely Pythonic.

Two main features of PyTorch are:

- Tensor computation with strong GPU acceleration support.

- Automatic Differentiation for creating and training deep neural networks.

Common PyTorch Module

In PyTorch, modules are used to represent neural networks.

i. Autograd

The autograd module is PyTorch's automatic differentiation that helps to compute the gradients in the forward pass quite time. Autograd generates a directed acyclic graph (DAG) to represent the input tensors while the nodes are the output

ii. Optim

The optim module is a package with pre-written algorithms for optimizers that can be used to build neural networks.

iii. nn

The nn module includes various classes that help to build neural network models. All modules in PyTorch inherit from nn module.

Conclusion: We learnt about the various deep learning

Assignment NO. 2

Name: Akhil X

Roll No: 004

Subject: LP-IV (DL)

DE-17 2022-23

DOF: 08 Aug 2022 DOC: 22 Aug 2022

Grade: 0/10 Sign: [Signature]

Aim: Implementing feedforward neural network with keras and tensorflow.

- Import the necessary packages.
- Load the training and testing data.
- Define the network architecture using keras.
- Train the model using SGD.
- Evaluate the network.
- Plot the training loss and accuracy.

Objective: To learn how to develop a feedforward neural network and how to optimize it for better performance.

Infrastructure: Computer/Laptop, Virtual Machine

Software used: Jupyter Notebook / Google Colab / JupyterLab and Keras

Theory:

What is feedforward neural network?

Deep feedforward networks, also often are called feedforward neural networks, or multilayer perceptrons. The goal of a feedforward neural network is to approximate some function $f: \mathbb{R}^n \rightarrow \mathbb{R}^m$. For example, for a classifier, $y = f(x)$ maps an input x to a category y . A feedforward network defines a mapping $y = f(x; \theta)$ and learns the value of the parameters θ that result in the best function approximation.


These models are called feedforward because information flows

through the function being evaluated from it, though it intermediate computations used to compute and finally output it. There are no feedback connections within the nodes are fed back into itself.

Structure of a feedforward Neural network

The basic structure consists of the following layers:

- Input layer: It is where the user accepts the input for the neural network.
- Hidden layer: This is the layer where all the computations required upon the prediction are done.
- Output layer: The output from the hidden layer is provided at the output layer.



The nodes are connected with the help of edges. The edge represented by 'w' is called 'weight'. It represents the node where the output is and represents the node where the edge ends.

The nodes compute the output for the next layer by the product of the node input and the weight across the node edge, which is then applied to an activation function.

decide whether the node should fire or not for the output.

SGD

Stochastic gradient descent (SGD) and its variants are probably the most used optimization algorithms for ML in general and for deep learning in particular. A crucial parameter for the SGD algorithm is the learning rate.

The standard gradient descent algorithm updates the parameters θ . If θ is the objective $J(\theta)$ then

$$\theta = \theta - \alpha \nabla_{\theta} J(\theta)$$

where the expectation in the above equation is approximated by evaluating the cost and gradient over the full training set.

MNIST / CIFAR 10

MNIST: The MNIST dataset of handwritten digit's has training set of 70,000 examples and each row of the matrix corresponding to a 28x28 image. The response variable is range from 0 to 9.

CIFAR 10: CIFAR-10 is an common establishment established computer vision dataset used for object recognition. The dataset have 60 million tiny images dataset and consist of 60,000, 32x32 color images containing one of 10 object classes. Further more, the data were converted from RGB to gray, normalized and standard to 2 decimal places.

Implementation

- Load the necessary libraries
- Load the datasets from the libraries or from internet.
- Build the feedforward neural network using keras.
- Train the model with the dataset and use SGD as optimizer.
- Evaluate the model for the accuracy and other evaluation metrics.
- Plot the loss and accuracy graph.

Conclusion: We developed a feedforward neural network for handwritten digit recognition.

Name: Akhil A

Roll No: 004


Subject: LP-IV(DL)

```
In [1]: import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import matplotlib.pyplot as plt
# Import seaborn as sns
```

MNIST dataset

```
In [2]: mnist = tf.keras.datasets.mnist
(x_train, y_train), (x_test, y_test) = mnist.load_data() # Data loading
x_train, x_test = x_train/255.0, x_test/255.0 # Normalizing the data
```

```
In [27]: sns.heatmap(x_train[0])
plt.show()
```



Preparing the model

```
In [3]: model = Sequential([
    Flatten(input_shape=(28, 28)),
    Dense(128, activation="relu"),
    Dropout(0.2),
    Dense(10)
])
```

```

In [5]: predictions = model(x_train[:1]),numpy()
Out[5]: array([[ -0.24707137, -0.64293617,  0.33793105, -0.7325163, -0.10029303,
  0.42578584, -0.5628654, -0.9137927, -1.2458755,  0.75219357]],
          dtype=float32)

In [6]: tf.nn.softmax(predictions).numpy()
Out[6]: array([[0.08687506, 0.0584754, 0.15438868, 0.05346493, 0.10060976,
  0.17026025, 0.06335013, 0.0446007, 0.03199779, 0.2397735]],
          dtype=float32)

In [8]: loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)

In [9]: model.compile(optimizer="adam", loss = loss_fn, metrics=['accuracy'])

In [10]: model.fit(x_train, y_train, epochs=5)

Epoch 1/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.2930 - accur
acy: 0.9143
Epoch 2/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1412 - accur
acy: 0.9575
Epoch 3/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.1048 - accur
acy: 0.9683
Epoch 4/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.0855 - accur
acy: 0.9736
Epoch 5/5
1875/1875 [=====] - 5s 3ms/step - loss: 0.0729 - accur
acy: 0.9769

Out[10]: <Keras.callbacks.History at 0x27c0bc71210>

In [11]: model.evaluate(x_test, y_test, verbose=2)
313/313 - 1s - loss: 0.0750 - accuracy: 0.9764 - 849ms/epoch - 3ms/step


Out[11]: [0.07503402578181839, 0.9764000177383423]

```

Validation of Model


```
In [12]: val = model.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test),
Epoch 1/5
300/300 [=====] - 1s 4ms/step - loss: 0.0516 - accurac
y: 0.9841 - val_loss: 0.0659 - val_accuracy: 0.9793
Epoch 2/5
300/300 [=====] - 1s 4ms/step - loss: 0.0456 - accurac
y: 0.9858 - val_loss: 0.0641 - val_accuracy: 0.9799
Epoch 3/5
300/300 [=====] - 1s 4ms/step - loss: 0.0437 - accurac
y: 0.9865 - val_loss: 0.0649 - val_accuracy: 0.9807
Epoch 4/5
300/300 [=====] - 1s 4ms/step - loss: 0.0415 - accurac
y: 0.9870 - val_loss: 0.0633 - val_accuracy: 0.9812
Epoch 5/5
300/300 [=====] - 1s 4ms/step - loss: 0.0398 - accurac
y: 0.9874 - val_loss: 0.0627 - val_accuracy: 0.9808
```

```
In [13]: plt.title("Model Accuracy")
plt.ylabel("Accuracy")
plt.xlabel("epoch")
plt.plot(val.history["accuracy"])
plt.plot(val.history["val_accuracy"])
plt.legend(["train", "val"])
plt.show()
```



In []:

ASSIGNMENT NO-3

Name : SHILPA

Roll No. : 003

Subject : LP-IV (DL)

GE - I 2022-23

DOB : 22 Aug 2022 DOB : 29 Aug 2022

Marked : 10/10 Sign : [Signature]

AIM: Build the image classification model by dividing the model into the following 4 stages:

- a. Loading the and preprocessing the image data.
- b. Defining the model architecture
- c. Training the model
- d. Estimating the model's performance.

Objective: To learn about CNN and how to develop a CNN for image recognition.

Infrastructure: Computer/Laptop/Virtual Machine

Software used: Jupyter Notebook/JupyterLab, TensorFlow, Keras

THEORY

CNN


Convolutional networks, also known as convolutional neural networks or CNNs, are a specialized kind of neural network for processing data that has a known, grid-like topology. The name "convolutional neural network" indicates that the network employs a mathematical operation called convolution. Convolution is a specialized kind of linear operation. Convolutional networks are simply neural networks that use convolution in place of general matrix multiplication in atleast one of their layers.

Architecture of CNN

- There are two main parts of a CNN:
- A convolution step that separates and identifies the various features

of the image for analysis is a process called as feature extraction. The network of feature extraction consists of many layers of convolutional or pooling layers.

- Fully connected layers that utilizes the output from the pooling layers process and predicts the class of the image or the features extracted in previous stages.
- This CNN model of feature extraction aims to reduce the of features present in a dataset. It creates new features & summarises the existing features contained in an image of pictures.



Convolution Layers

These are three types of layers that make up the CNN which are 1. convolutional layers, 2. pooling layers, and 3. fully-connected (FC) layers. In addition to these layers there are three more layers important parameters which are the dropout layer and the activation function.

i. Convolutional Layers

This layer is the first layer that is used to extract

features from the input image. In this layer, the mathematical operation of cross-correlation is performed between the input image and a filter of a particular size $m \times n$. By sliding the filter over the input image, the dot product is taken between the filter and the pixels of the input image with respect to the size of the filter (CNN).

The output is derived as the feature map which gives us information about the image such as the corners and edges.

ii. Pooling layers

The primary aim of this layer is to decrease the size of the convolved feature map to reduce the computational cost. This is performed by decreasing the correlation between layers and independently operates on each feature map. In Max Pooling, the largest element is taken from feature map. Average pooling calculates the average of the elements in a predefined sized image section. The total sum of the elements in the predefined section is computed in down pooling.

iii. Fully Connected Layer

The fully connected (FC) layer consists of the weights and biases along with the neurons and is used to connect the neurons between two different layers. These layers are usually placed before the output layer and form the last few layers of a CNN architecture.

iv. Dropout

Usually, when all the features are connected to the FC layer, it can cause overfitting in the training dataset. Overfitting occurs when a particular

Title of E

model is able to on the training data causing a negative impact on the model's performance when used on new data.

V. Abstraction Function

They are used to learn and approximate any kind of local and complex relationship between variables of the network in simple words, it decides which information of the network should give in the forward direction and which one should not at the end of the network.

Implementation

- Load the necessary libraries.
- Import the dataset from the respective library on local.
- Design the neural network architecture and mention the layers, nodes, etc.
- Train the model with the imported dataset.
- Evaluate the performance of the model.

Conclusion :

We learnt how to build and train a CNN to identify images.

Name:- Akhil A
Roll No:- 04
PRN No:- 72036135E
Class:- BE(II)
Subject:- Deep Learning

```
In [1]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
import numpy as np
import random
import matplotlib.pyplot as plt
import tensorflow as tf
```

```
In [2]: import os
os.environ['CUDA_VISIBLE_DEVICES'] = '-1'

Load Dataset
```

Dataset available on <https://tf.keras.github.io/tfjs-backend-cuda/>

```
In [3]: x_train = np.loadtxt('input.csv', delimiter=',')
y_train = np.loadtxt('labels.csv', delimiter=',')

x_test = np.loadtxt('input_test.csv', delimiter=',')
y_test = np.loadtxt('labels_test.csv', delimiter=',')
```

```
In [4]: x_train = x_train.reshape(len(x_train), 100, 100, 3)
y_train = y_train.reshape(len(y_train), 1)

x_test = x_test.reshape(len(x_test), 100, 100, 3)
y_test = y_test.reshape(len(y_test), 1)

x_train = x_train/255.0
x_test = x_test/255.0
```

```
In [5]: print("Shape of X_train:", x_train.shape)
print("Shape of y_train:", y_train.shape)
print("Shape of X_test:", x_test.shape)
print("Shape of y_test:", y_test.shape)

Shape of X_train: (2000, 100, 100, 3)
Shape of y_train: (2000, 1)
Shape of X_test: (400, 100, 100, 3)
Shape of y_test: (400, 1)
```

```
In [6]: idx = Random.randint(0, len(x_train))
plt.imshow(x_train[idx,:])
plt.show()
```



Model Building

```
In [8]: model = Sequential([
    Conv2D(256,(3,3),activation = 'relu',input_shape=(100,100,3)),
    BatchNormalization(),
    MaxPooling2D((4,4)),
    Conv2D(128,(3,3),activation = 'relu'),
    BatchNormalization(),
    MaxPooling2D((2,2)),
    Conv2D(64,(3,3),activation = 'relu'),
    BatchNormalization(),
    MaxPooling2D((2,2)),
    Flatten(),
    Dense(128,activation='relu'),
    Dropout(0.4),
    Dense(1,activation='sigmoid')
])
```

```
In [9]: opt = SGD(momentum=0.9)
model.compile(optimizer=opt, loss='binary_crossentropy', metrics = ['accuracy'])
```



```

In [10]: model.fit(x_train, y_train, epochs=30, steps_per_epoch = 20, validation_data=(x_test,
Epoch 1/30
20/20 [=====] - 61s 3s/step - loss: 0.8490 - accuracy:
0.5650 - val_loss: 0.6948 - val_accuracy: 0.5250
Epoch 2/30
20/20 [=====] - 70s 4s/step - loss: 0.6686 - accuracy:
0.6240 - val_loss: 0.6801 - val_accuracy: 0.5975
Epoch 3/30
20/20 [=====] - 68s 3s/step - loss: 0.6057 - accuracy:
0.6830 - val_loss: 0.6797 - val_accuracy: 0.5625
Epoch 4/30
20/20 [=====] - 64s 3s/step - loss: 0.5674 - accuracy:
0.6990 - val_loss: 0.7254 - val_accuracy: 0.5050
Epoch 5/30
20/20 [=====] - 64s 3s/step - loss: 0.5089 - accuracy:
0.7605 - val_loss: 0.7999 - val_accuracy: 0.5050
Epoch 6/30
20/20 [=====] - 62s 3s/step - loss: 0.4879 - accuracy:
0.7720 - val_loss: 0.8164 - val_accuracy: 0.4950
Epoch 7/30
20/20 [=====] - 61s 3s/step - loss: 0.4430 - accuracy:
0.7960 - val_loss: 0.8590 - val_accuracy: 0.5125
Epoch 8/30
20/20 [=====] - 60s 3s/step - loss: 0.4141 - accuracy:
0.8130 - val_loss: 0.8978 - val_accuracy: 0.5050
Epoch 9/30
20/20 [=====] - 65s 3s/step - loss: 0.3859 - accuracy:
0.8270 - val_loss: 1.0060 - val_accuracy: 0.5025
Epoch 10/30
20/20 [=====] - 59s 3s/step - loss: 0.3405 - accuracy:
0.8470 - val_loss: 1.0907 - val_accuracy: 0.5200
Epoch 11/30
20/20 [=====] - 59s 3s/step - loss: 0.3248 - accuracy:
0.8555 - val_loss: 0.9742 - val_accuracy: 0.5300
Epoch 12/30
20/20 [=====] - 59s 3s/step - loss: 0.2908 - accuracy:
0.8820 - val_loss: 0.8911 - val_accuracy: 0.5650
Epoch 13/30
20/20 [=====] - 59s 3s/step - loss: 0.2615 - accuracy:
0.8920 - val_loss: 0.6965 - val_accuracy: 0.6225
Epoch 14/30
20/20 [=====] - 59s 3s/step - loss: 0.2234 - accuracy:
0.9145 - val_loss: 0.7834 - val_accuracy: 0.6000
Epoch 15/30
20/20 [=====] - 60s 3s/step - loss: 0.1961 - accuracy:
0.9245 - val_loss: 1.3311 - val_accuracy: 0.5475
Epoch 16/30
20/20 [=====] - 60s 3s/step - loss: 0.1657 - accuracy:
0.9470 - val_loss: 1.0894 - val_accuracy: 0.5900
Epoch 17/30
20/20 [=====] - 60s 3s/step - loss: 0.1325 - accuracy:
0.9590 - val_loss: 1.3612 - val_accuracy: 0.5525
Epoch 18/30
20/20 [=====] - 60s 3s/step - loss: 0.1255 - accuracy:
0.9565 - val_loss: 1.0850 - val_accuracy: 0.6050
Epoch 19/30

```

```

20/20 [=====] - 60s 3s/step - loss: 0.1015 - accuracy:
Epoch 20/30
20/20 [=====] - 60s 3s/step - loss: 0.0923 - accuracy:
Epoch 21/30
20/20 [=====] - 60s 3s/step - loss: 0.0923 - accuracy:
Epoch 22/30
20/20 [=====] - 60s 3s/step - loss: 0.0725 - accuracy:
Epoch 23/30
20/20 [=====] - 60s 3s/step - loss: 0.0725 - accuracy:
Epoch 24/30
20/20 [=====] - 60s 3s/step - loss: 0.1011 - accuracy:
Epoch 25/30
20/20 [=====] - 60s 3s/step - loss: 0.0831 - accuracy:
Epoch 26/30
20/20 [=====] - 60s 3s/step - loss: 0.0670 - accuracy:
Epoch 27/30
20/20 [=====] - 60s 3s/step - loss: 0.0483 - accuracy:
Epoch 28/30
20/20 [=====] - 60s 3s/step - loss: 0.0448 - accuracy:
Epoch 29/30
20/20 [=====] - 60s 3s/step - loss: 0.0401 - accuracy:
Epoch 30/30
20/20 [=====] - 60s 3s/step - loss: 0.0376 - accuracy:
Epoch 31/30
20/20 [=====] - 60s 3s/step - loss: 0.0272 - accuracy:
Epoch 32/30
20/20 [=====] - 60s 3s/step - loss: 0.0204 - accuracy:

```

```
Out[10]: <keras.callbacks.History at 0x1c39aeace00>
```

```
In [11]: model.evaluate(x_test,y_test)
```

```
13/13 [=====] - 2s 164ms/step - loss: 1.0204 - accuracy:
y: 0.7100
```

```
Out[11]: [1.0204460620880127, 0.709999785423279]
```

Making Predictions



```

In [14]: idx2 = random.randint(0, len(y_test))
plt.imshow(x_test[idx2,:])
plt.show()

y_pred = model.predict(x_test[idx2,:].reshape(1,100,100,3))
y_pred = y_pred > 0.5
if (y_pred==0):
    pred = 'dog'
else:
    pred = 'cat'

print("Our model says it is a", pred)

```



```

1/1 [=====] - 0s 28ms/step
Our model says it is a dog

```

```
In [15]: score = model.evaluate(X_test, Y_test, verbose = 0)
print("Test Score: ", score[0])
print("Test accuracy: ", score[1])

Test Score: 1.0204460620886127
Test accuracy: 0.709999785423279
```


```
In [16]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 98, 98, 256)	7168
batch_normalization (Batch Normalization)	(None, 98, 98, 256)	1824
max_pooling2d (MaxPooling2D)	(None, 24, 24, 256)	0
conv2d_1 (Conv2D)	(None, 22, 22, 128)	295040
batch_normalization_1 (Batch Normalization)	(None, 22, 22, 128)	512
max_pooling2d_1 (MaxPooling2D)	(None, 11, 11, 128)	0
conv2d_2 (Conv2D)	(None, 9, 9, 64)	73792
batch_normalization_2 (Batch Normalization)	(None, 9, 9, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 4, 4, 64)	0
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 128)	131200
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129
Total params: 509,121		
Trainable params: 508,225		
Non-trainable params: 896		

```
In [17]: val = model.fit(x_train, y_train, epochs=5, validation_data=(x_test, y_test), batch_
10/10 [=====] - 61s 6s/step - loss: 0.0246 - accuracy:
0.9925 - val_loss: 0.9134 - val_accuracy: 0.7250
Epoch 2/5
10/10 [=====] - 58s 6s/step - loss: 0.0173 - accuracy:
0.9970 - val_loss: 0.8590 - val_accuracy: 0.7275
Epoch 3/5
10/10 [=====] - 58s 6s/step - loss: 0.0139 - accuracy:
0.9970 - val_loss: 0.8365 - val_accuracy: 0.7425
Epoch 4/5
10/10 [=====] - 59s 6s/step - loss: 0.0101 - accuracy:
0.9995 - val_loss: 0.8777 - val_accuracy: 0.7425
Epoch 5/5
10/10 [=====] - 58s 6s/step - loss: 0.0116 - accuracy:
0.9975 - val_loss: 0.8869 - val_accuracy: 0.7650
```

```
In [18]: plt.title("Model Accuracy")
plt.xlabel("Accuracy")
plt.ylabel("epoch")
plt.plot(val.history['accuracy'])
plt.plot(val.history['val_accuracy'])
plt.legend(['train', 'val'])
plt.show()
```



Name: XHJL X

Roll No: 004

Assignment No: 4

SUBJECT : L.P-IV (DL)

2022-23

NBM SDE - 21

DOP: 29 Aug 2022 DOC: 12 Sept 2022

Mark: 10/10 Sign: Pr.

STM: Use Autoencoders to implement anomaly detection. Build the model by using

1. Dataset required for training.
2. Upload / access the dataset.
3. Encoder convert it into latent representation.
4. Decoder reconstructs it back to the original input.
5. Compare the model with training loss and Evaluation Metrics.

OBJECTIVE: To learn about Autoencoders and developing autoencoders for anomaly detection.

Infrastructure: Computer / Laptop / Virtual Machine

Software used: Jupyter Notebook / Google Colab, TensorFlow, Keras

THEORY


AUTO ENCODER

Autoencoders are artificial neural network capable of learning efficient representation of the input data, called encoding, without using supervision. These coding have typically a much lower dimensionality than the input data, making autoencoders useful for dimensionality reduction and compression. These coding's, the code is a compact "summary" or "compression" of the input.

Autoencoders act as powerful feature detectors and can be used for unsupervised pre-training of deep neural networks. Similarly, they can be capable of randomly generating new data that looks very similar to the training data. For example, you can train an encoder

on picture of faces and it would then be able to generate faces.

Architecture of Autoencoders



An autoencoder consists of 3 components:

- i. Encoder: It compresses the input into a latent space representation. The encoder layer encodes the input image as a compressed representation in a reduced dimension. The compressed image is the distributed version of the original image.
- ii. Code: This is the compressed input (from encoder) which is used to decode for reconstructing the original input later.
- iii. Decoder: It decodes the encoded output in form of code to original input. The decoded output is a latent reconstruction of original input. The goal is to get an output as close as the input.

The danger between the encoder and decoder i.e. the code known as bottleneck. This is a well-designed approach to

which aspect of the observed data are relevant information and which aspects can be discarded.

Parameters used for training an autoencoder:

- From parameters are:
 - i. Code size: It is a number of nodes in the middle layer. smaller size results in more compression and it may be difficult to make the size smaller beyond a certain limit to get satisfactory results.
 - ii. Number of layers: The autoencoders can be as deep as you like. They are very similar to an ANN, you only need to decide how many layers autoencoders should have.
 - iii. Number of nodes per layer: Number of nodes per layer decreases with each subsequent layer of the encoder and increases back in the decoder. Also the decoder is usually symmetric to the encoder in terms of layer structure.
 - iv. Loss function: We use either MSE (Mean Square Error) or Binary cross entropy as the loss function. If the input values are in the range [0,1] then you typically use cross-entropy, otherwise you use the mean square error.

Implementation:

- i. Load the necessary libraries.
- ii. Import the dataset from the respective library.
- iii. Shape the data as per your needs.
- iv. Encode the input data in latent representation.
- v. Decode the output of the encoder to connect it back to original input.
- vi. Use the models with Optimization, loss, and Evaluation metrics.

Conclusion:
We learn how to detect anomaly using ensemble-subspace

Name:- Akhil A
Roll No:- 04
PRN No:- 7206135E
Class:- BE(II)
Subject:- Deep Learning

```
In [1]: import matplotlib.pyplot as plt
import pandas as pd
import tensorflow as tf
import seaborn as sns
from tensorflow.keras.models import Model
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.model_selection import train_test_split
sns.set()
import numpy as np
from tensorflow.keras import layers, losses
```

```
In [2]: df = pd.read_csv('http://storage.googleapis.com/download.tensorflow.org/data/ecg.
raw_data = df.values
df.head()
```

```
Out[2]:
```

	0	1	2	3	4	5	6	7	8
0	-0.112522	-2.827204	-3.773897	-4.348751	-4.376041	-3.474986	-2.181408	-1.818286	-1.250522
1	-1.100878	-3.996840	-4.285843	-4.506579	-4.022377	-3.234368	-1.566126	-0.092258	-0.754680
2	-0.567088	-2.593450	-3.874230	-4.584096	-4.187449	-3.151462	-1.742940	-1.490609	-1.183980
3	0.490473	-1.914407	-3.616364	-4.318823	-4.268016	-3.881110	-2.993280	-1.871131	-1.333884
4	0.800232	-0.874252	-2.384761	-3.973292	-4.338224	-3.802422	-2.534510	-1.783423	-1.094450

5 rows x 141 columns

```
In [3]: labels = raw_data[:, -1]
data = raw_data[:, 0:-1]
```

```
In [4]: pd.Series(labels).value_counts()
```

```
Out[4]: 1.0    2919
0.0    2079
dtype: int64
```

```
In [5]: train_data, test_data, train_labels, test_labels = train_test_split(
data, labels, test_size = 0.2, random_state=21
)
```

Title of I

Study of

Application

Image C

Abstract

Keywords

Goal of the

```
In [6]: min = np.min(train_data)
         max = np.max(train_data)
         train_data = ( train_data - min ) / ( max - min )
         test_data = ( test_data - min ) / ( max - min )
```

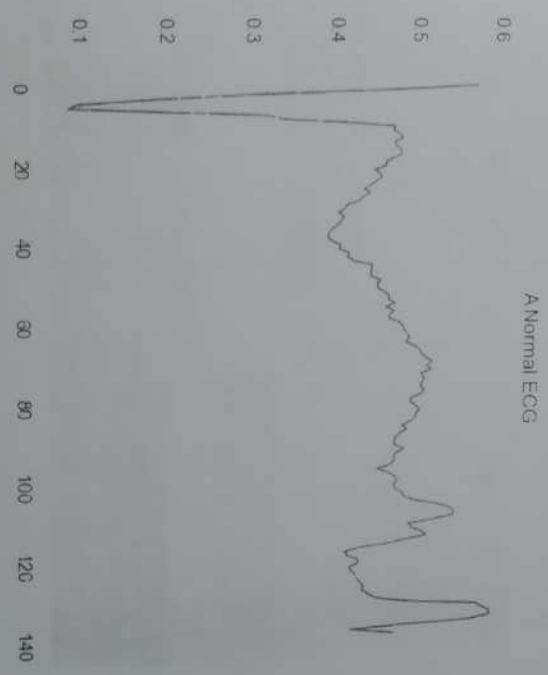
```
In [7]: train_labels = train_labels.astype(bool)
         test_labels = test_labels.astype(bool)
```

```
normal_train_data = train_data[train_labels]
normal_test_data = test_data[test_labels]
anomalous_train_data = train_data[~train_labels]
anomalous_test_data = test_data[~test_labels]
```


```
In [8]: ~train_labels
```

```
Out[8]: array([False, False, False, ..., False, False, False])
```

```
In [9]: plt.grid()
         plt.plot(np.arange(140), normal_train_data[9])
         plt.title('A Normal ECG')
         plt.show()
```




```
In [10]: plt.grid()
plt.plot(np.arange(140),anomalous_train_data[0])
plt.title('A Anomalous ECG')
plt.show()
```



```
In [11]: class AnomalyDetector(Model):
def __init__(self):
    super(AnomalyDetector, self).__init__()
    self.encoder = Sequential([
        layers.Dense(32, activation='relu'),
        layers.Dense(16, activation='relu'),
        layers.Dense(8, activation='relu')
    ])
    self.decoder = tf.keras.Sequential([
        layers.Dense(16, activation='relu'),
        layers.Dense(32, activation='relu'),
        layers.Dense(140, activation='sigmoid')
    ])

def call(self, x):
    encoded = self.encoder(x)
    decoded = self.decoder(encoded)
    return decoded

autoencoder = AnomalyDetector()
```

```
In [12]: autoencoder.compile(optimizer='adam', loss='mse')
```

```

In [13]: history = autoencoder.fit(normal_train_data, normal_train_data,
                                   epochs = 20,
                                   batch_size=312,
                                   validation_data=(normal_test_data, normal_test_data),
                                   shuffle=True)

Epoch 1/20
5/5 [=====] - 1s 45ms/step - loss: 0.0598 - val_loss:
0.0565
Epoch 2/20
5/5 [=====] - 0s 10ms/step - loss: 0.0559 - val_loss:
0.0545
Epoch 3/20
5/5 [=====] - 0s 10ms/step - loss: 0.0533 - val_loss:
0.0508
Epoch 4/20
5/5 [=====] - 0s 10ms/step - loss: 0.0494 - val_loss:
0.0470
Epoch 5/20
5/5 [=====] - 0s 9ms/step - loss: 0.0456 - val_loss:
0.0428
Epoch 6/20
5/5 [=====] - 0s 10ms/step - loss: 0.0413 - val_loss:
0.0387
Epoch 7/20
5/5 [=====] - 0s 9ms/step - loss: 0.0376 - val_loss:
0.0354
Epoch 8/20
5/5 [=====] - 0s 10ms/step - loss: 0.0344 - val_loss:
0.0326
Epoch 9/20
5/5 [=====] - 0s 10ms/step - loss: 0.0319 - val_loss:
0.0303
Epoch 10/20
5/5 [=====] - 0s 10ms/step - loss: 0.0297 - val_loss:
0.0283
Epoch 11/20
5/5 [=====] - 0s 10ms/step - loss: 0.0279 - val_loss:
0.0267
Epoch 12/20
5/5 [=====] - 0s 9ms/step - loss: 0.0264 - val_loss:
0.0253
Epoch 13/20
5/5 [=====] - 0s 10ms/step - loss: 0.0251 - val_loss:
0.0243
Epoch 14/20
5/5 [=====] - 0s 10ms/step - loss: 0.0242 - val_loss:
0.0235
Epoch 15/20
5/5 [=====] - 0s 10ms/step - loss: 0.0235 - val_loss:
0.0229
Epoch 16/20
5/5 [=====] - 0s 10ms/step - loss: 0.0230 - val_loss:
0.0224
Epoch 17/20
5/5 [=====] - 0s 10ms/step - loss: 0.0225 - val_loss:
0.0219


```



```
Epoch 18/20
5/5 [=====] - 0s 9ms/step - loss: 0.0210 - val_loss:
0.0214
Epoch 19/20
5/5 [=====] - 0s 10ms/step - loss: 0.0215 - val_loss:
0.0209
Epoch 20/20
5/5 [=====] - 0s 9ms/step - loss: 0.0210 - val_loss:
0.0206
```

```
In [14]: plt.plot(history.history['loss'],label='Training Loss')
plt.plot(history.history['val_loss'],label='Testing Loss')
plt.legend()
```

```
Out[14]: <matplotlib.legend at 0x23e12fe7790>
```




```

In [15]: encoded_image = autoencoder.encoder(normal_test_data).numpy()
          decoded_image = autoencoder.decoder(encoded_image).numpy()

          plt.plot(normal_test_data[0], 'b')
          plt.plot(decoded_image[0], 'r')
          plt.fill_between(np.arange(148), decoded_image[0], normal_test_data[0], color='lightgray')
          plt.legend(labels=['Input', 'Reconstruction', 'Error'])
          plt.show()


```



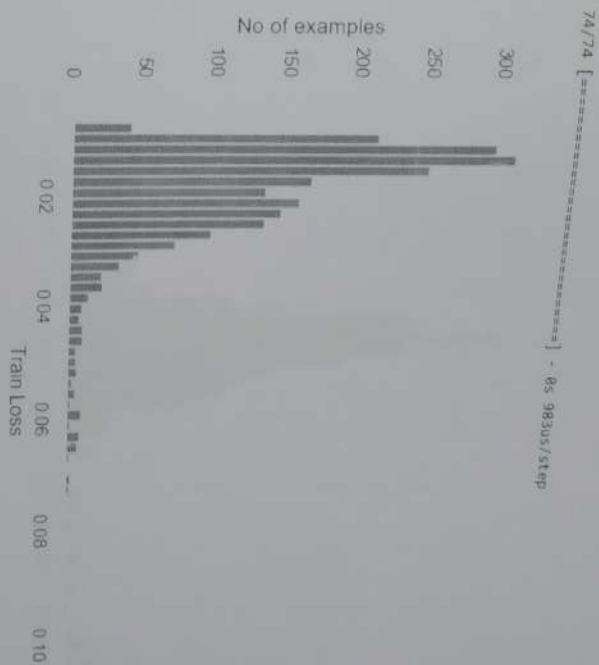
```

In [16]: encoded_image = autoencoder.encoder(anamalous_test_data).numpy()
         decoded_image = autoencoder.decoder(encoded_image).numpy()
         plt.plot(anamalous_test_data[0], 'b')
         plt.plot(decoded_image[0], 'r')
         plt.fill_between(np.arange(148), decoded_image[0], anamalous_test_data[0], color='g')
         plt.legend(labels=['Input', 'Reconstruction', 'Error'])
         plt.show()

```



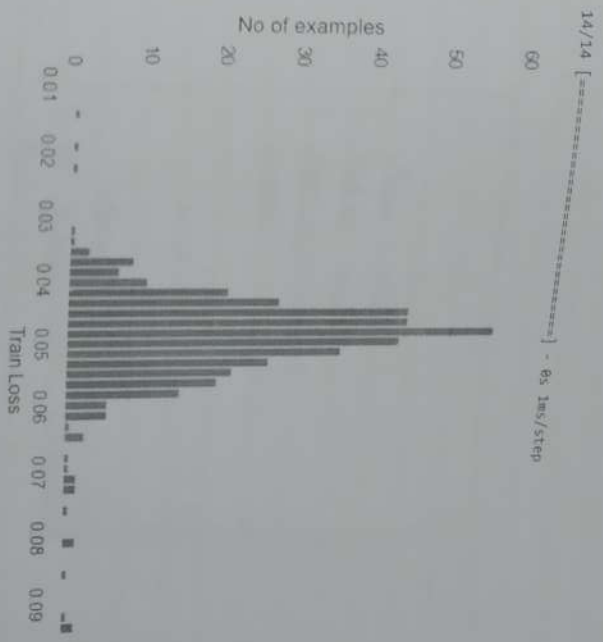

```
In [17]: reconstructions = autoencoder.predict(normal_train_data)
train_loss = tf.keras.losses.mae(reconstructions, normal_train_data)
plt.hist(train_loss[None,:], bins=50)
plt.xlabel("Train Loss")
plt.ylabel("No of examples")
plt.show()
```



```
In [18]: threshold = np.mean(train_loss) + np.std(train_loss)
print("Threshold: ", threshold)

Threshold: 0.03308283181023525
```

```
In [19]: reconstructions = autoencoder.predict(anomalous_test_data)
test_loss = tf.keras.losses.mae(reconstructions, anomalous_test_data)
plt.hist(test_loss[None,:], bins=50)
plt.xlabel("Train Loss")
plt.ylabel("No of examples")
plt.show()
```



```
In [20]: def predict(model,data,threshold):
reconstructions = model(data)
loss = tf.keras.losses.mae(reconstructions,data)
return tf.math.less(loss, threshold)

def print_stats(predictions, labels):
print("Accuracy = {}".format(accuracy_score(labels,preds)))
print("Precision = {}".format(precision_score(labels,preds)))
print("Recall = {}".format(recall_score(labels,preds)))
```

```
In [21]: preds = predict(autoencoder, test_data, threshold)
print_stats(preds, test_labels)

Accuracy = 0.945
Precision = 0.99287298448343
Recall = 0.9089285714285714
```

Name : MHEL X

Roll No : 004

Assignment NO : 5

SUBJECT : LP-IV (DL)

2022-23

NSNBOE-11

DDP : 12 Aug 2022 Doc : 19 Aug 2022

Mark : 10/10 Sign : 

AIM: Implement the Longshorts Bag of Words (LGBOW) Model.

- Data preparation.
- Generate training data.
- Train model.
- Output.

OBJECTIVE: To learn and understand Longshorts bag of words model.

Infrastructure: Computer / Laptop / Virtual Machine

Software used: TensorFlow / Keras / PyTorch, Jupyter, etc.

THEORY:

What is the LGBOW model?

The LGBOW model tries to understand the context of the words and takes this as input. It then tries to predict words that are contextually accurate. Let us consider an example for understanding. Consider the sentence: "It is a pleasant day." and the word "pleasant" goes as an input to the neural network. We are trying to predict the word 'day' here. We will use the one-hot encoding for the input words and measure the error ratio with the one-hot encoded target word. Doing this will help us predict the output with based on the word with least error.

The model architecture


- The LGBOW model architecture tries to predict the target word by trying to understand the context of the surrounding words.

Consider if the sentence as above, "It is a pleasant day".
model converts this sentence into words pairs in the form
(context word, target word). The user will have to set the hidden
size. If the window for the context word is 2 then the
pairs would look like this: $([it, a], [is])$, $([is, pleasant], [a])$,
 $([pleasant], [day])$. With these word pairs, the model tries to predict the
word considered the context words.

If we have 4 context words used for predicting one target
word the input layer will be in the form of four 1×4 input
vectors. These input vectors will be passed to the hidden layer
entire the same layers where an element-wise summation
performed on the vectors & before a final activation is performed
and the output is obtained.

Implementation of the CBOW model

- Import the libraries and read our dataset.
- For the implementation of this model, we will use a sample
test data.
- Generate function that create window size and pairs of
words
- Build neural network on sample data.



Conclusion:

We saw what is CBOW model and how it works. There can be
used for text recognition, speech to text conversion, etc.

Name:- Akhil A
Roll No:- 84
PRN No:- 72036133E
Class:- BE[IT]
Subject:- DL

```
In [2]: import numpy as np
import tensorflow as tf
import keras.backend as K
from tensorflow.keras.models import Sequential
from keras.layers import Dense, Embedding, Lambda
from keras.preprocessing import text, sequence
import gensim
```

```
In [4]: data=open('covid.txt','r')
corona_data = [text for text in data if text.count(' ') >= 2]
vectorize = Tokenizer()
vectorize.fit_on_texts(corona_data)
corona_data = vectorize.texts_to_sequences(corona_data)
total_vocab = sum(len(s) for s in corona_data)
word_count = len(vectorize.word_index) + 1
window_size = 2
```

```
In [11]: def chow_model(data, window_size, total_vocab):
total_length = window_size*2
for text in data:
text_len = len(text)
for idx, word in enumerate(text):
context_word = []
target = []
begin = idx - window_size
end = idx + window_size + 1
context_word.append(text[i] for i in range(begin, end) if 0 <= i < 1)
target.append(word)
contextual = sequence.pad_sequences(context_word, total_length=total)
final_target = np_utils.to_categorical(target, total_vocab)
yield(contextual, final_target)
```


In [15]:

```
model = Sequential()
model.add(Embedding(input_dim=total_vocab, output_dim=100, input_len=
model.add(Dense(input_dim=x: K.mean(x, axis=1), output_shape=(100,)))
for i in range(20):
    cost = 0
    for x, y in crow_model(data, window_size, total_vocab):
        cost += model.train_on_batch(contextual, final_target)
    print(i, cost)
```

```
0 0
1 0
2 0
3 0
4 0
5 0
6 0
7 0
8 0
9 0
10 0
11 0
12 0
13 0
14 0
15 0
16 0
17 0
18 0
19 0
```

In [16]: dimensions=100

```
vect_file = open('vectors.txt', 'w')
vect_file.write('{} {} \n'.format(100, dimensions))
```

Out[16]: 8

```

In [17]: weights = model.get_weights()[0]
for text, i in vectorize_word_index:
    final_vec = ''.join(map(str, list(weights[i, :])))
    vect_file.write('{} {}{}\n'.format(text, final_vec))
    vect_file.close()

```

```

cbow_output = gensim.models.KeyedVector.load_word2vec_format('vectors.txt', bins=
cbow_output.most_similar(positive=['virus'])

```

```

Out[17]: [(['generated', 0.2327983673872375),
('is', 0.22727039456367493),
('present', 0.1759325639076233),
('appear', 0.1655784398317337),
('not', 0.164816975935669),
('between', 0.1457787901631012),
('days', 0.1345373988151588),
('the', 0.1339661031961441),
('does', 0.11900028586387634),
('pre', 0.11821701377630234)]

```

```

In [ ]:

```

Name: AMIT K

Roll No: 004

Assignment No: 6

Subject: LP-IV (DL)

2022-23

NBNSIDE-11

DOP: 14 Sep 2022 Doc: 03 Oct 2022

Mark: (10) Sign *[Signature]*

Step: Object detection using Transfer Learning

1. Load in a pre-trained CNN model trained on a large dataset.
2. Freeze parameters (weights) in model's lower convolutional layers.
3. Add custom classifier with several layers of trainable parameters to model.
4. Train classifier layers on training data available for task.
5. Fine-tune higher parameters and unfreeze more layers as needed.

Objective: To load a pre-trained model and improve its performance by Transfer Learning architecture.

Infrastructure: Computer / Laptop / Virtual Machine

Software used: TensorFlow / Keras / Google Colab, Tensorflow, Keras

THEORY

What is Transfer Learning?

Transfer Learning generally refers to a process where a model trained on one problem is used in some way on a second related problem. One or more layers from the trained model are used \neq in a new model trained on the problem of interest.

Transfer Learning has the benefit of decreasing the training time for a neural network model and can result in lower generalization error.

The weights in re-used layers may be used as the starting point for the training process and adapted in response to the new problem. This way helps transfer learning as a type of weight

indistinguishable. This may be useful when the first related problem has a lot more labelled data than the problem of interest and the similarity in the structure of the problem may be useful in both contexts.

How to use pre-trained models:

Some of these usage patterns as follows:

- i. Classifier: The pre-trained model is used directly to classify new images.
- ii. Discriminative Feature Extractor: The pre-trained model, or some part of the model, is used to pre-process images and extract relevant features.
- iii. Integrated Feature Extractor: The pre-trained model, or some part of the model, is integrated into a new model, but largely of a pre-trained model are frozen during training.
- iv. Weight Initialization: The pre-trained, or some part of the model, is integrated into a new model, and the layers of the pre-trained model are trained in concert with the new model.

It may not be clear as to which usage of the pre-trained model may yield the best results on your new computer vision task. Therefore some experimentation may be required.

Ways to Fine-tune the model

- i. Feature Extraction: We can use a pre-trained model as a feature extraction mechanism. What we can do is that we can remove the top layers and then use the entire network as a fixed feature extractor for the new data set.

extractor for the new data set.

We use the architecture of the pre-trained model. What we can do is that we use architecture of the model while we initialize all the weights randomly and train the model according to our dataset again. Train some layers while freeze others. Another way to use a pre-trained model is to train it partially. What we can do is that we keep the weights of initial layers of the model frozen while we retrain only the higher layers. We can try and test as to how many layers to be frozen and how many to be trained.

Building a Deep Learning-Based Object Detection Model

Training for a performing deep learning model for object detection takes a lot of data and computing power. To facilitate the development, we can use transfer learning by using training and testing pre-trained based on other related datasets.

Since there are multiple backend libraries such as PyTorch and Hyperparameter to take care of when training a full-scale deep learning model, we can create a central dictionary to store these configuration parameters, including setting up the different paths, installing relevant libraries, and downloading the pre-trained models.

Conclusion:

We concluded from the experiment, how to develop a model for a specific application with the help of transfer learning architecture in deep learning.

Assignment No 6

November 9, 2022

Name: Akhil A Roll No.: 004 PRN: 72036135E Subject: LP-IV(DL)

```

import tensorflow as tf
import numpy as np
import cv2
import PIL.Image as Image
import os
import matplotlib.pyplot as plt
import tensorflow_hub as hub
import pathlib

```

```

img_shape = (224, 224)

```

```

url_dataset = "https://storage.googleapis.com/download.tensorflow.org/
example_images/flower_photos.tgz"

```

```

data_dir = tf.keras.utils.get_file(origin=url_dataset,
                                     filename='flower_photos',
                                     untar=True)
data_dir = pathlib.Path(data_dir)

```

```

Downloading data from https://storage.googleapis.com/download.tensorflow.org/exa
mple_images/flower_photos.tgz
22813984/22813984 [=====] - 1s 0us/step

```

```

# Total images
image_count = len(list(data_dir.glob('*/*.jpg')))
print(image_count)

```

3670

0.0.1 Making the different classes

```

flowers_images_dict = {
    "daisy": list(data_dir.glob('daisy/*')),
    "dandelion": list(data_dir.glob('dandelion/*')),
    "roses": list(data_dir.glob('roses/*')),
    "sunflowers": list(data_dir.glob('sunflowers/*')),
    "tulips": list(data_dir.glob('tulips/*'))
}

```



```

}

flowers_labels_dict = {
    "daisy": 0,
    "dandelion": 1,
    "roses": 2,
    "sunflowers": 3,
    "tulips": 4
}

```

0.0.2 Reshaping Images

```

g) X, Y = [], []

for flower_name, images in flowers_images_dict.items():
    for image in images:
        img = cv2.imread(str(image))
        resized_img = cv2.resize(img, Image.Shape)
        X.append(resized_img)
        Y.append(flowers_labels_dict[flower_name])

d) X = np.array(X)
    Y = np.array(Y)

e) from sklearn.model_selection import train_test_split
    X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

f) X_train_scaled = X_train / 255
    X_test_scaled = X_test / 255

```

0.0.3 Pretrained Model

```

h) tf_model = "https://tfhub.dev/google/tf2-preview/mobilenet_v2/feature_vector/4"

i) classifier = tf.keras.Sequential([
    hub.KerasLayer(tf_model, input_shape=(224, 224, 3), trainable=False),
    tf.keras.layers.Dense(len(flowers_labels_dict), activation="softmax")
])

classifier.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
keras_layer (KerasLayer)	(None, 1280)	2257984
dense (Dense)	(None, 5)	6405

```

=====
total params: 2,264,389
trainable params: 6,405
non-trainable params: 2,257,984
=====

```

```

4: classifier.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
    metrics=['accuracy'])

```

```

5: classifier.fit(X_train_scaled, y_train, epochs=5)
Epoch 1/5

```

```

/usr/local/lib/python3.7/dist-packages/tensorflow/python/util/dispatch.py:1082:
UserWarning: "sparse_categorical_crossentropy" received "from_logits=True", but
the "output" argument was produced by a sigmoid or softmax activation and thus
does not represent logits. Was this intended?
return dispatch_target(*args, **kwargs)

```

```

86/86 [=====] - 15s 42ms/step - loss: 0.8199 -
accuracy: 0.6962
Epoch 2/5

```

```

86/86 [=====] - 4s 45ms/step - loss: 0.4254 - accuracy:
0.8503
Epoch 3/5

```

```

86/86 [=====] - 4s 46ms/step - loss: 0.3299 - accuracy:
0.8950
Epoch 4/5

```

```

86/86 [=====] - 4s 43ms/step - loss: 0.2745 - accuracy:
0.9179
Epoch 5/5

```

```

86/86 [=====] - 4s 47ms/step - loss: 0.2363 - accuracy:
0.9320

```

```

<keras.callbacks.History at 0x7f76500871d0>

```

```

: classifier.evaluate(X_test_scaled, y_test)

```

```

29/29 [=====] - 2s 55ms/step - loss: 0.3484 - accuracy:
0.8725

```

```

[0.348428338766098, 0.8725489974021912]

```

0.0.4 Prediction

```
from PIL import Image
```

```
img = Image.open("360_F_105573812_cvd4f5j06tpp2Uk324q0YFb0pX1s0.jpg")
```

```
img
```



```
img = tf.keras.preprocessing.image.img_to_array(img, resize(Image.Shape))
```

```
img = np.array([img])
```

```
res = classifier.predict(img)
```

```
1/1 [=====] - 1s 737ms/step
```

```
print("The prediction is : {}".format(list(flowers_labels_dict.keys()[np.argmax(res)]))
```

```
The prediction is : roses
```