

```

In [5]: # example of using a pre-trained model as a classifier
from tensorflow.keras.preprocessing.image import load_img
from tensorflow.keras.preprocessing.image import img_to_array
from keras.applications.vgg16 import preprocess_input
from keras.applications.vgg16 import decode_predictions
from keras.applications.vgg16 import VGG16

In [19]: # load an image from file
image = load_img('download.jpg', target_size=(224, 224))

In [20]: # convert the image pixels to a numpy array
image = img_to_array(image)

In [21]: # reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))

In [22]: # prepare the image for the VGG model
image = preprocess_input(image)

In [23]: # load the model
model = VGG16()

In [24]: # predict the probability across all output classes
yhat = model.predict(image)

WARNING:tensorflow:5 out of the last 5 calls to <function Model.make_predict_function.<locals>.predict_function
at 0x000001F2498996C0> triggered tf.function retracing. Tracing is expensive and the excessive number of tracing
s could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3)
passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2)
, @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to h
ttps://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf
/function for more details.
1/1 [=====] - 1s 726ms/step

In [25]: # convert the probabilities to class labels
label = decode_predictions(yhat)
# convert the probabilities to class labels
label = decode_predictions(yhat)

In [26]: # retrieve the most likely result, e.g. highest probability
label = label[0][0]

In [27]: # print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))

castle (34.03%)

In [28]: # load an image from file
image = load_img('download2.png', target_size=(224, 224))

In [29]: # convert the image pixels to a numpy array
image = img_to_array(image)

In [30]: # reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))

In [31]: # prepare the image for the VGG model
image = preprocess_input(image)

In [32]: # load the model
model = VGG16()

In [33]: # predict the probability across all output classes
yhat = model.predict(image)

WARNING:tensorflow:6 out of the last 6 calls to <function Model.make_predict_function.<locals>.predict_function
at 0x000001F249898860> triggered tf.function retracing. Tracing is expensive and the excessive number of tracing
s could be due to (1) creating @tf.function repeatedly in a loop, (2) passing tensors with different shapes, (3)
passing Python objects instead of tensors. For (1), please define your @tf.function outside of the loop. For (2)
, @tf.function has reduce_retracing=True option that can avoid unnecessary retracing. For (3), please refer to h
ttps://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf
/function for more details.
1/1 [=====] - 1s 1000ms/step

In [34]: # convert the probabilities to class labels
label = decode_predictions(yhat)

In [35]: # retrieve the most likely result, e.g. highest probability
label = label[0][0]

```

```
label
```

```
Out[35]: ('n09468604', 'valley', 0.448454)
```

```
In [36]: # print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))

valley (44.85%)
```

```
In [37]: # load an image from file
image = load_img('download3.jpg', target_size=(224, 224))
```

```
In [38]: # convert the image pixels to a numpy array
image = img_to_array(image)
```

```
In [39]: # reshape data for the model
image = image.reshape((1, image.shape[0], image.shape[1], image.shape[2]))
```

```
In [46]: # prepare the image for the VGG model
image = preprocess_input(image)
```

```
In [47]: # load the model
model = VGG16()
```

```
In [48]: # predict the probability across all output classes
yhat = model.predict(image)
```

```
1/1 [=====] - 2s 2s/step
```

```
In [49]: # convert the probabilities to class labels
label = decode_predictions(yhat)
```

```
In [50]: # retrieve the most likely result, e.g. highest probability
label = label[0][0]
```

```
In [51]: # print the classification
print('%s (%.2f%%)' % (label[1], label[2]*100))

golden_retriever (77.67%)
```

```
In [52]: model.summary()
```

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_8 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
fc1 (Dense)	(None, 4096)	102764544
fc2 (Dense)	(None, 4096)	16781312
predictions (Dense)	(None, 1000)	4097000
=====		
Total params: 138357544 (527.79 MB)		
Trainable params: 138357544 (527.79 MB)		
Non-trainable params: 0 (0.00 Byte)		

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js